# Practical approach to sizing a Kubernetes cluster

- If total k8s node memory is approx. 96 GB, how many pods can be safely supported on it?
- Assuming all pods are equi-sized and requiring min(request) = 1 GB and max(limit) = 4 GB
- Not all pods will require 4 GB memory simulataneously
- Assuming, reliable gc occurring in the applications within the pods, at any given point of time, one can expect the average memory per pod to be 2.5 GB. So we can expect 96/2.5=38.4 pods to be supported, going by averages
- Max number of pods would be 96/1 = 96 pods
- Min number of pods would be 96/4 = 24 pods
- The number pods supported could range from 24 (low risk) to 96 (high risk of instability)
- Going by averages we should allow 38.4 pods, but, if we study the min-max memory requirements of our application, we can do a better job at predicting our apps “steady state” memory requirements per pod, say 2 GB and plan for 96/2 = 48 pods

- Total pods: 20
- All pods same size
- vCpu: min & max
- Memory: min & max

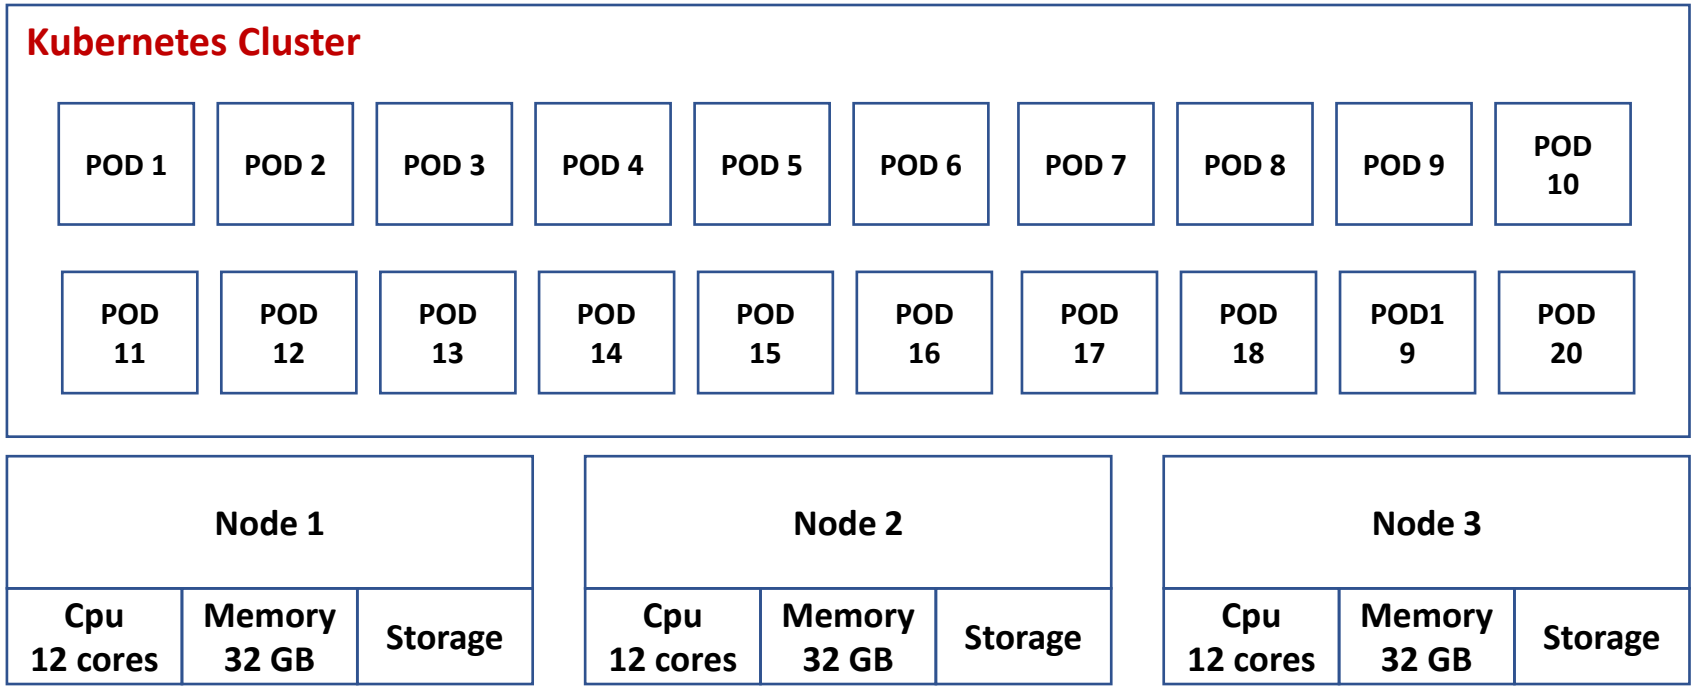


- Total cpu cores: 36
- Total Memory: 96 GB

# Practical RPA

- Often, most time consuming activity is data entry
- Automating data capture, using RPA & API calls
- **Leap Frog** UI development, with API automation

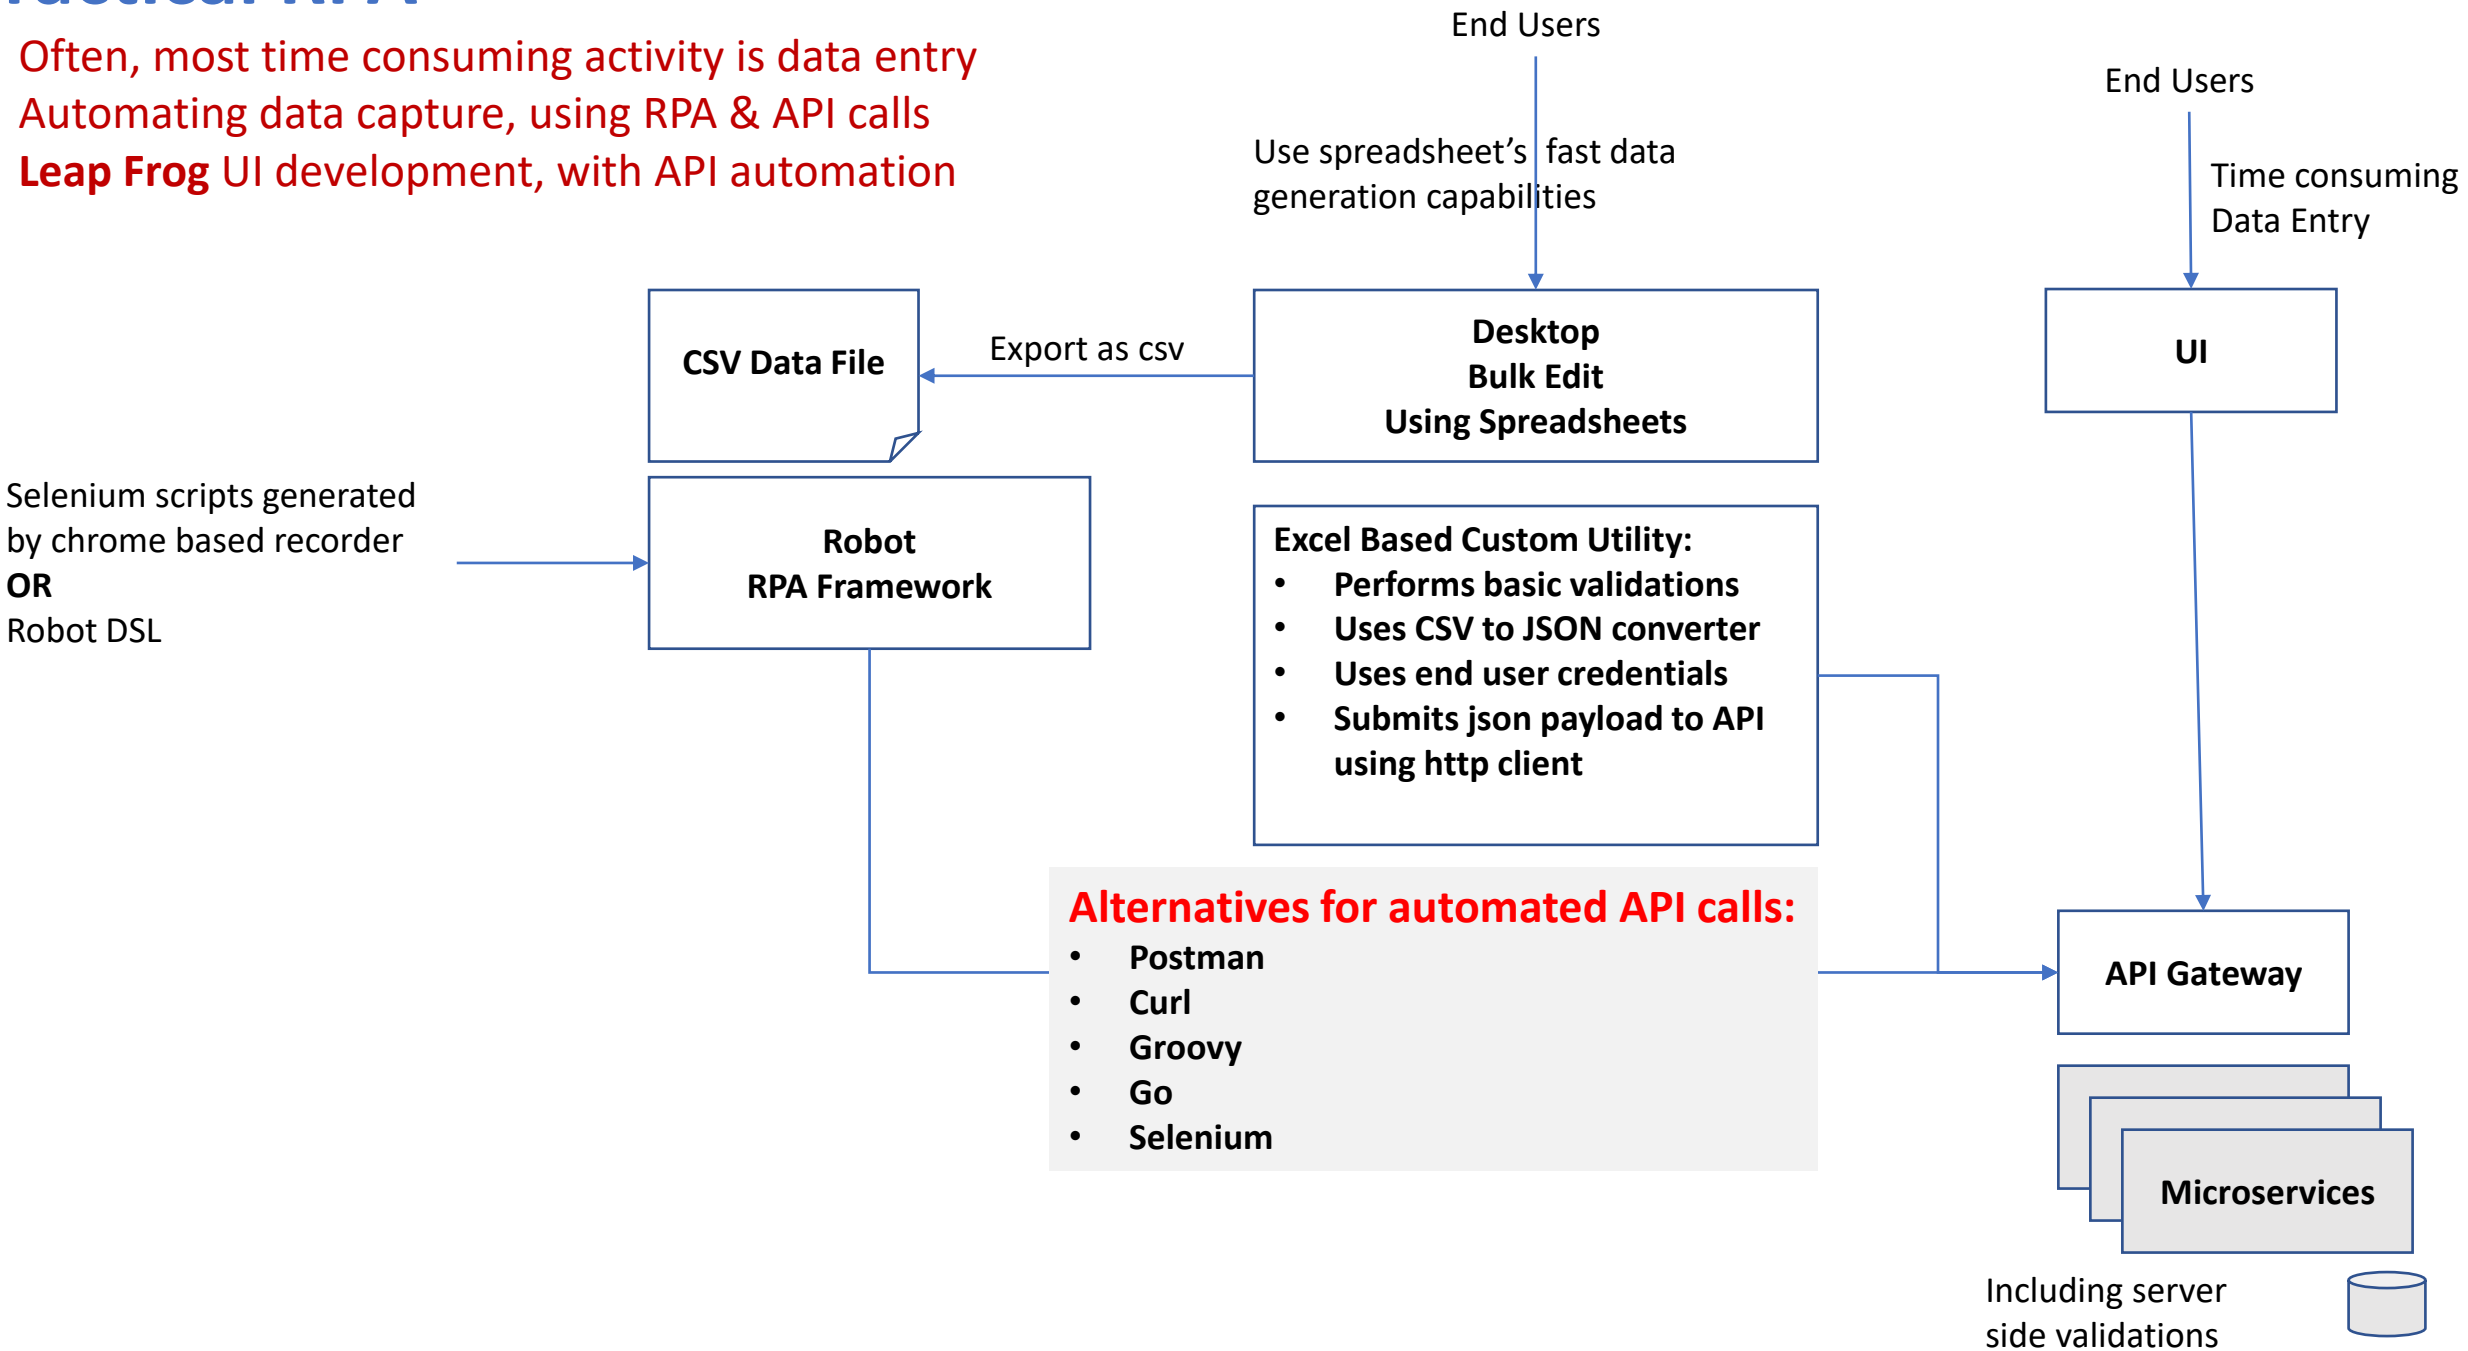

Cloud based low code development platforms
Eg. PowerApps, Outsystems, Mendix

**Enterprise Requirements from Low Code providers:**

- Secure access from low code servers to enterprise GW
- Secure Auth integrated/SSO with enterprise IAM
- Tenant isolation of access and data in SAAS cloud
  - Dedicated low code servers and access per enterprise
- VPC network extending from SAAS servers to on prem network

## Enterprise Apps Deployments

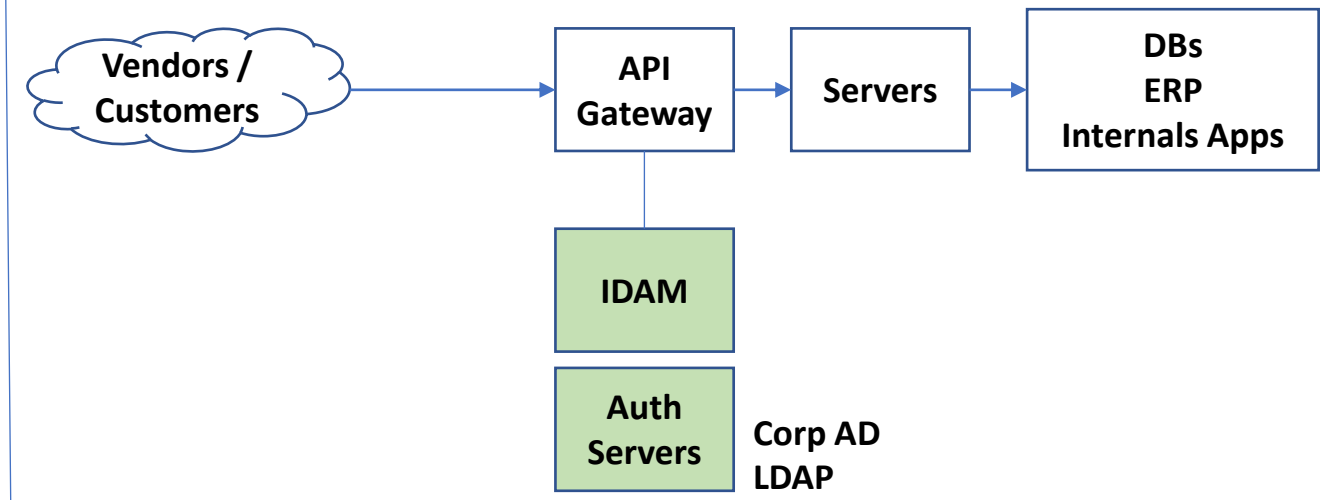


## Low Code based Enterprise Apps Deployments

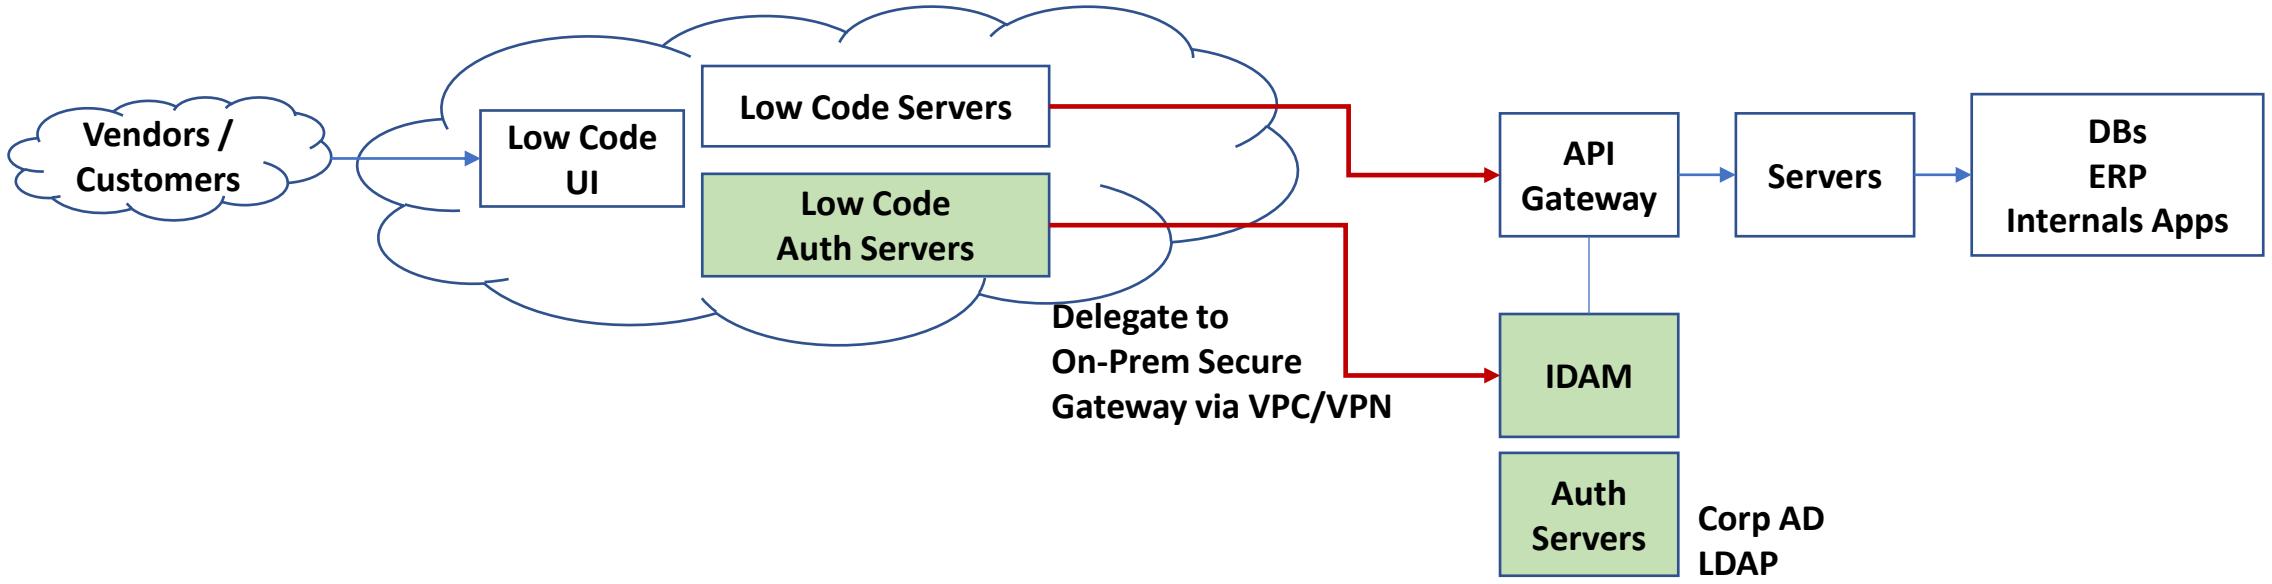

# Traditional API Implementation

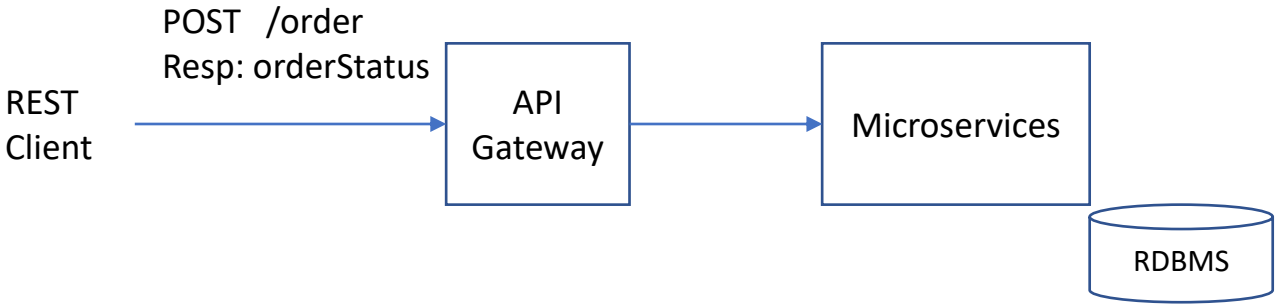


# High Scalability API Implementation

**Aysnc + Scale Out, all the way upto persistence**

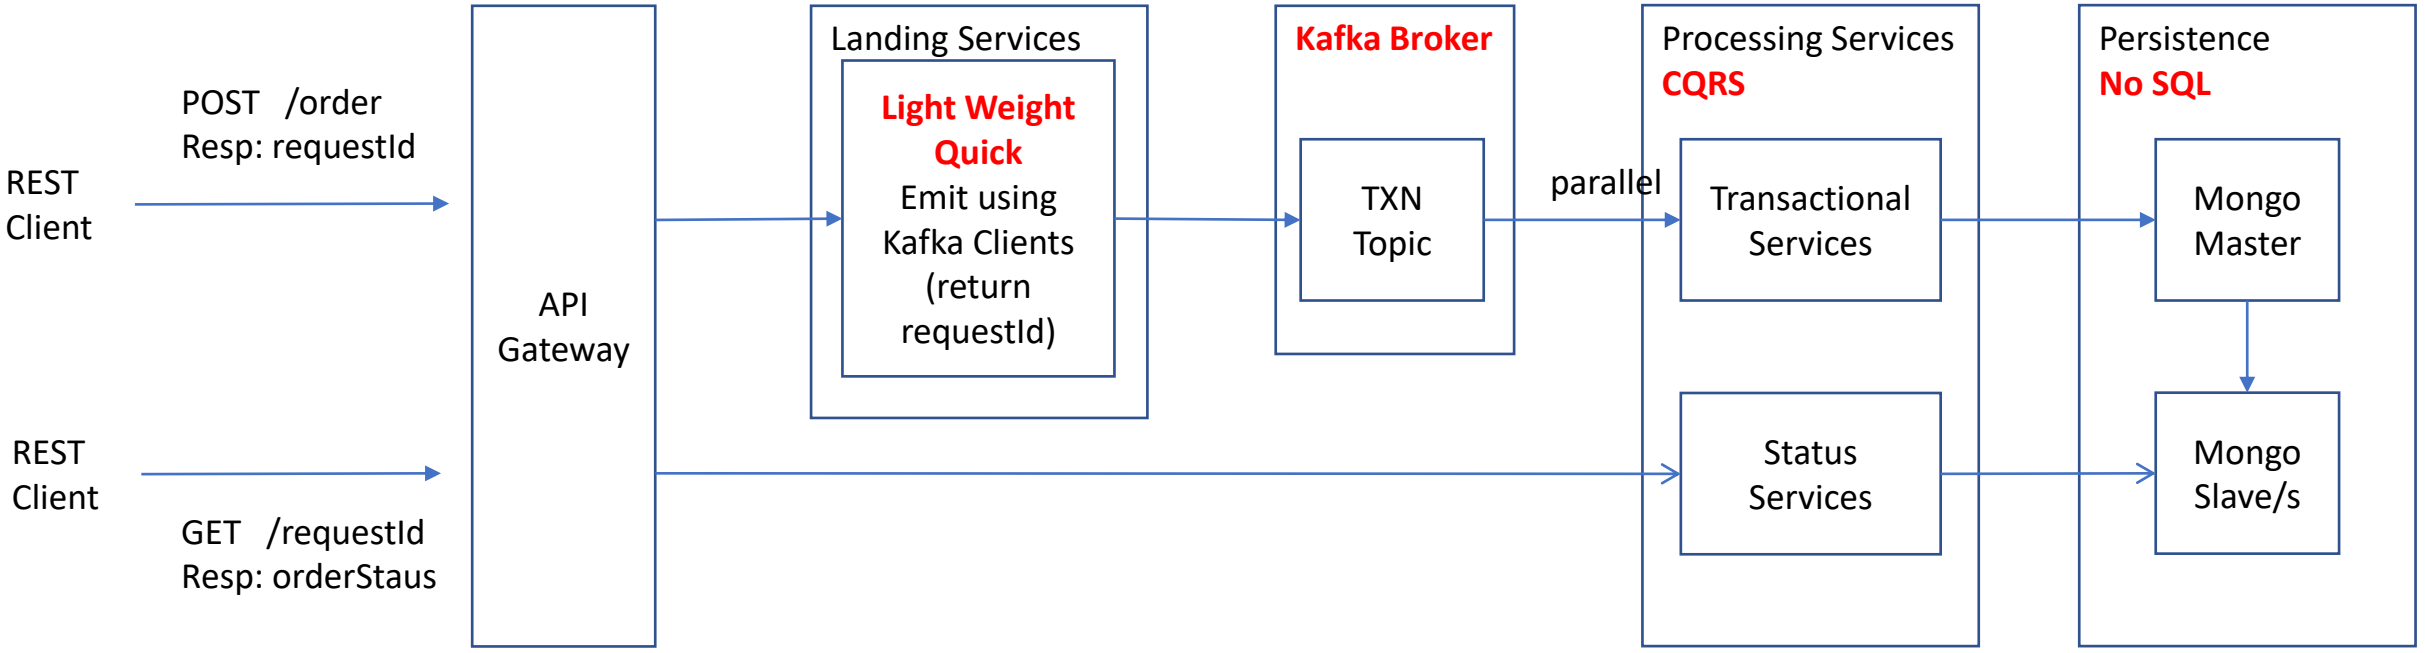

- Don’t forcibly create bounded contexts where none exist
- Instead hive out NFR services

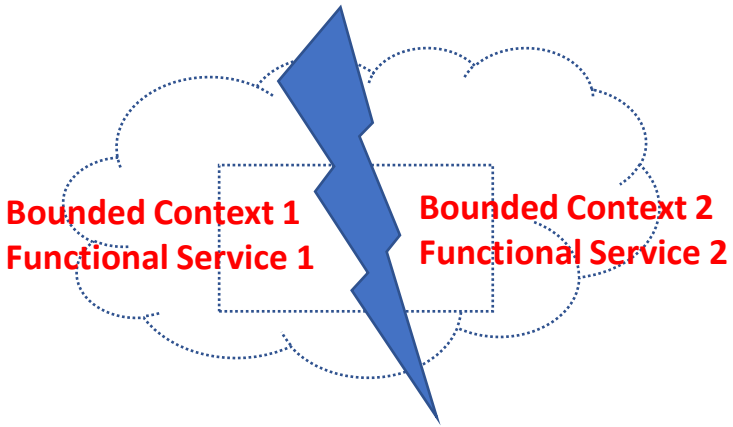


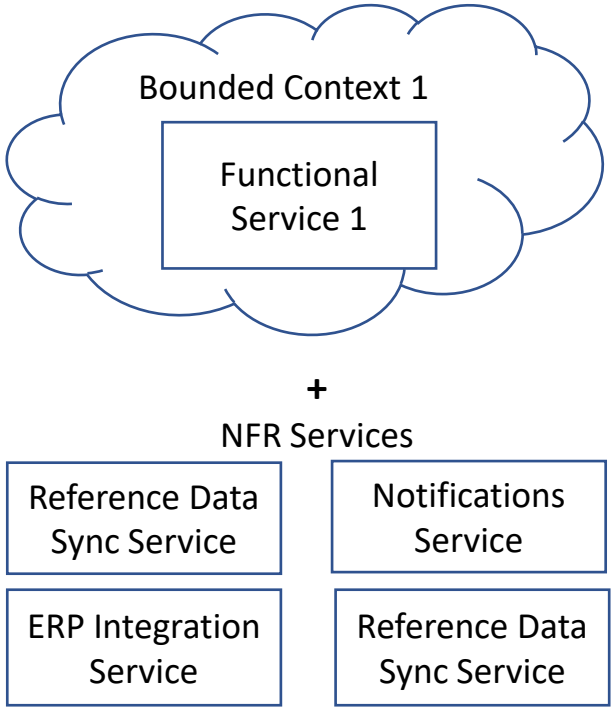


- To reduce http traffic / chattiness between microservices, consolidate API calls, where it makes sense

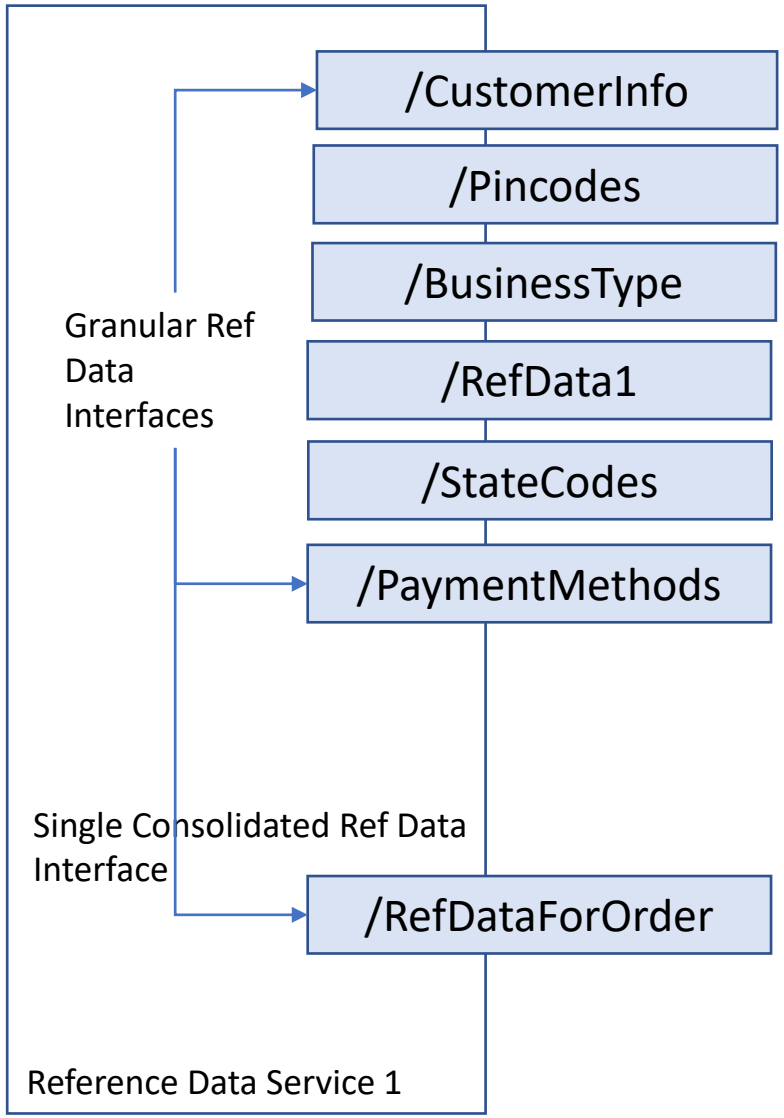


- Distinct Data model / microservice

**Ideal**

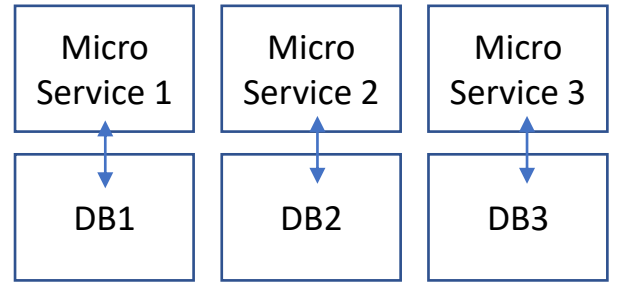


**Less than ideal**

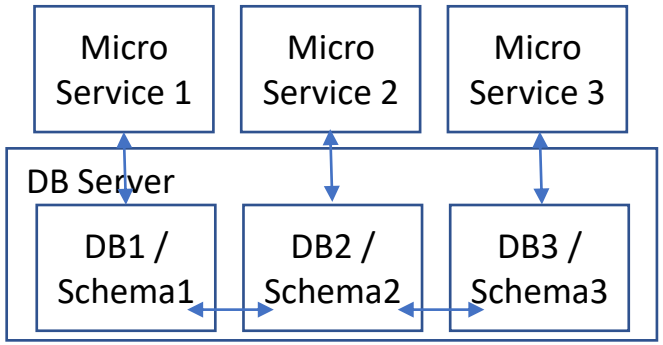


**Rubbish, with microservices**

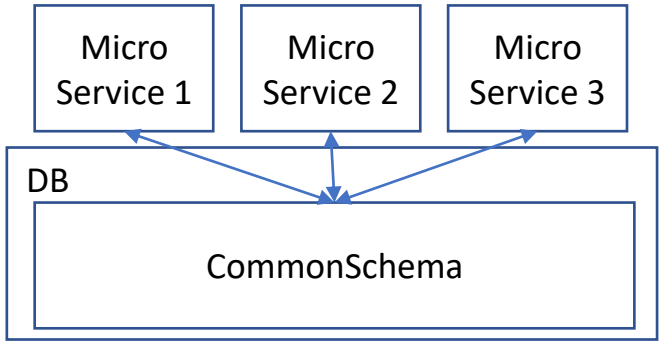

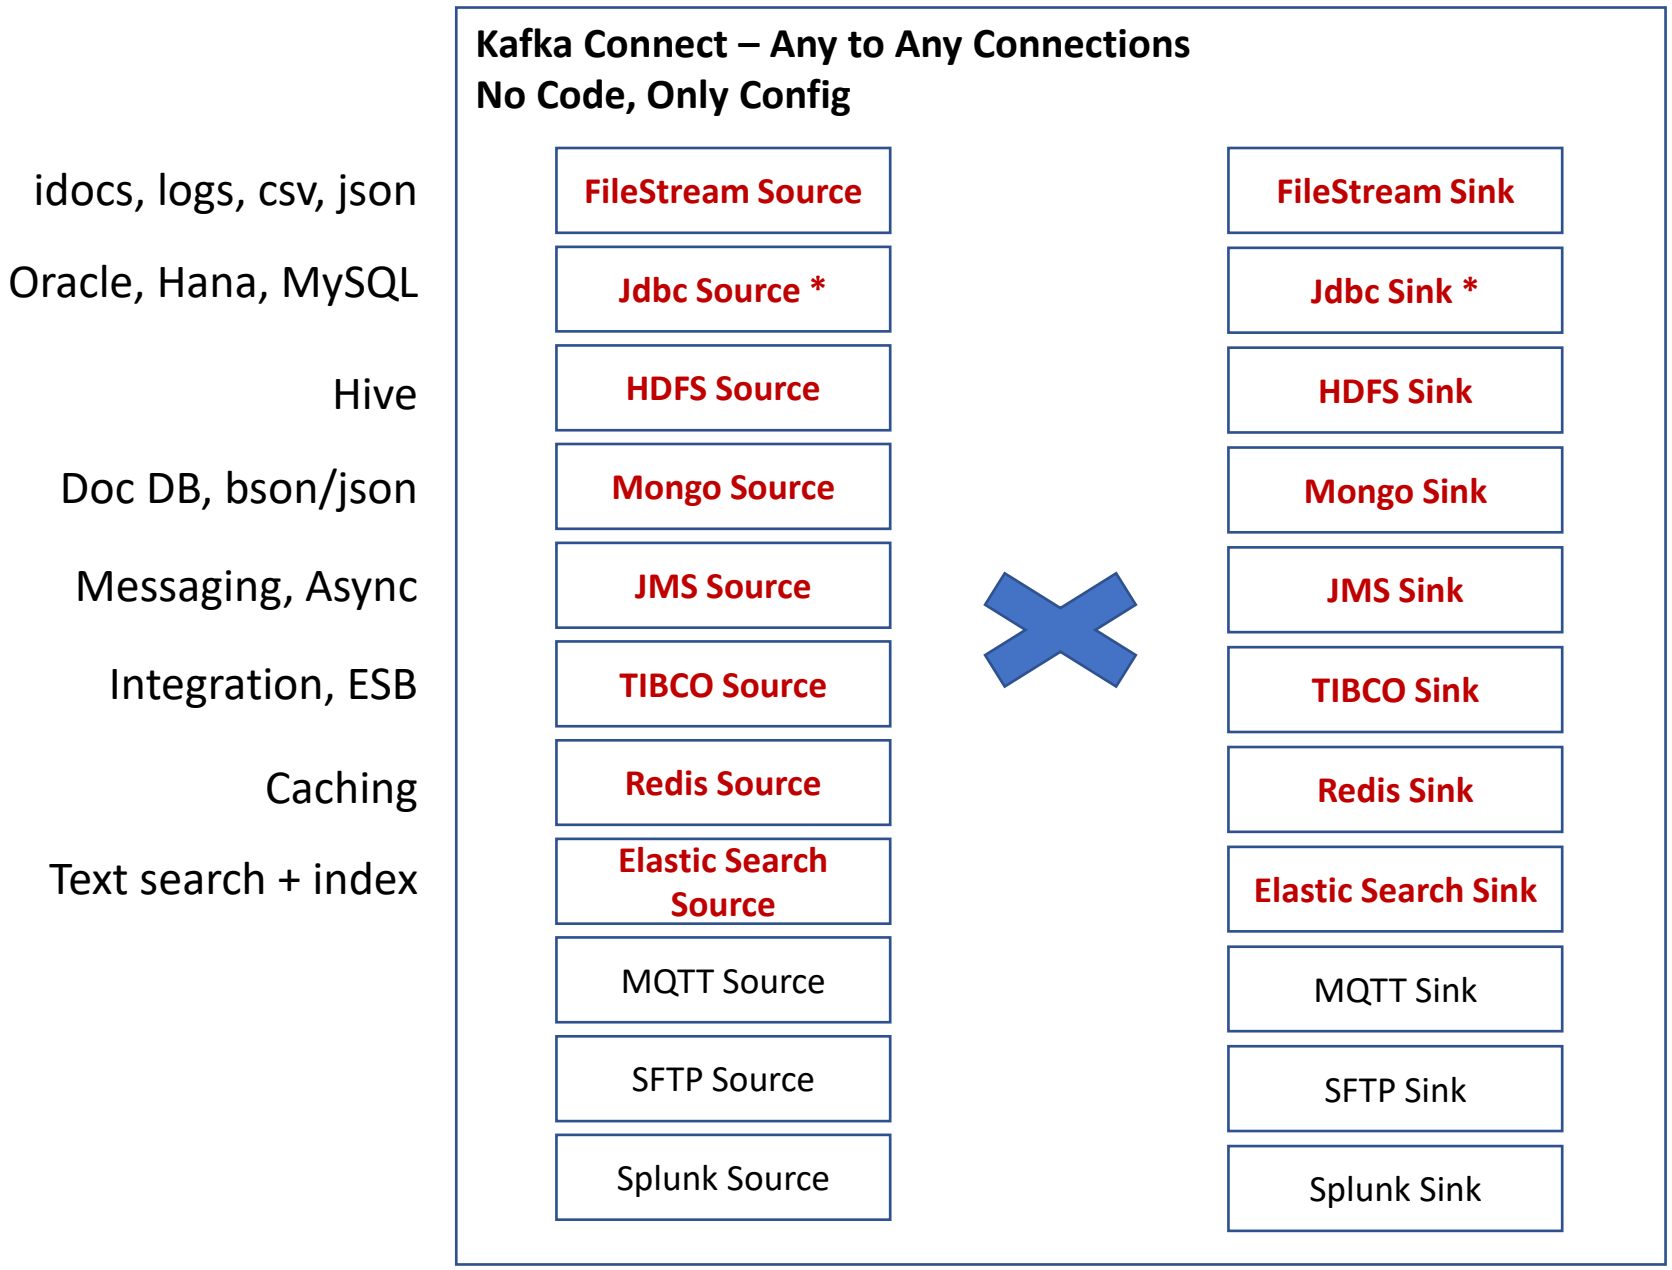
Kafka Connect – Any to Any Connections
No Code, Only Config
idocs, logs, csv, json
Oracle, Hana, MySQL
Hive
Doc DB, bson/json
Messaging, Async
Integration, ESB
Caching
Text search + index
FileStream Source
Jdbc Source *
HDFS Source
Mongo Source
JMS Source
TIBCO Source
Redis Source
Elastic Search Source
MQTT Source
SFTP Source
Splunk Source
FileStream Sink
Jdbc Sink *
HDFS Sink
Mongo Sink
JMS Sink
TIBCO Sink
Redis Sink
Elastic Search Sink
MQTT Sink
SFTP Sink
Splunk Sink
* = schema mandatory

# High throughput, batch processing, using Kafka and Mongo DB

**Partner Integration based on bulk or event based incoming data**

**Leverage Kafka Connectors or write custom kafka producers**

**Leverage Kafka IO Scalability for the landing platform**

**Throttle Consumption as per your resources**

**Use Kafka based async processing for better throughput (TPS)**

**Mongo DB for scalable persistence**

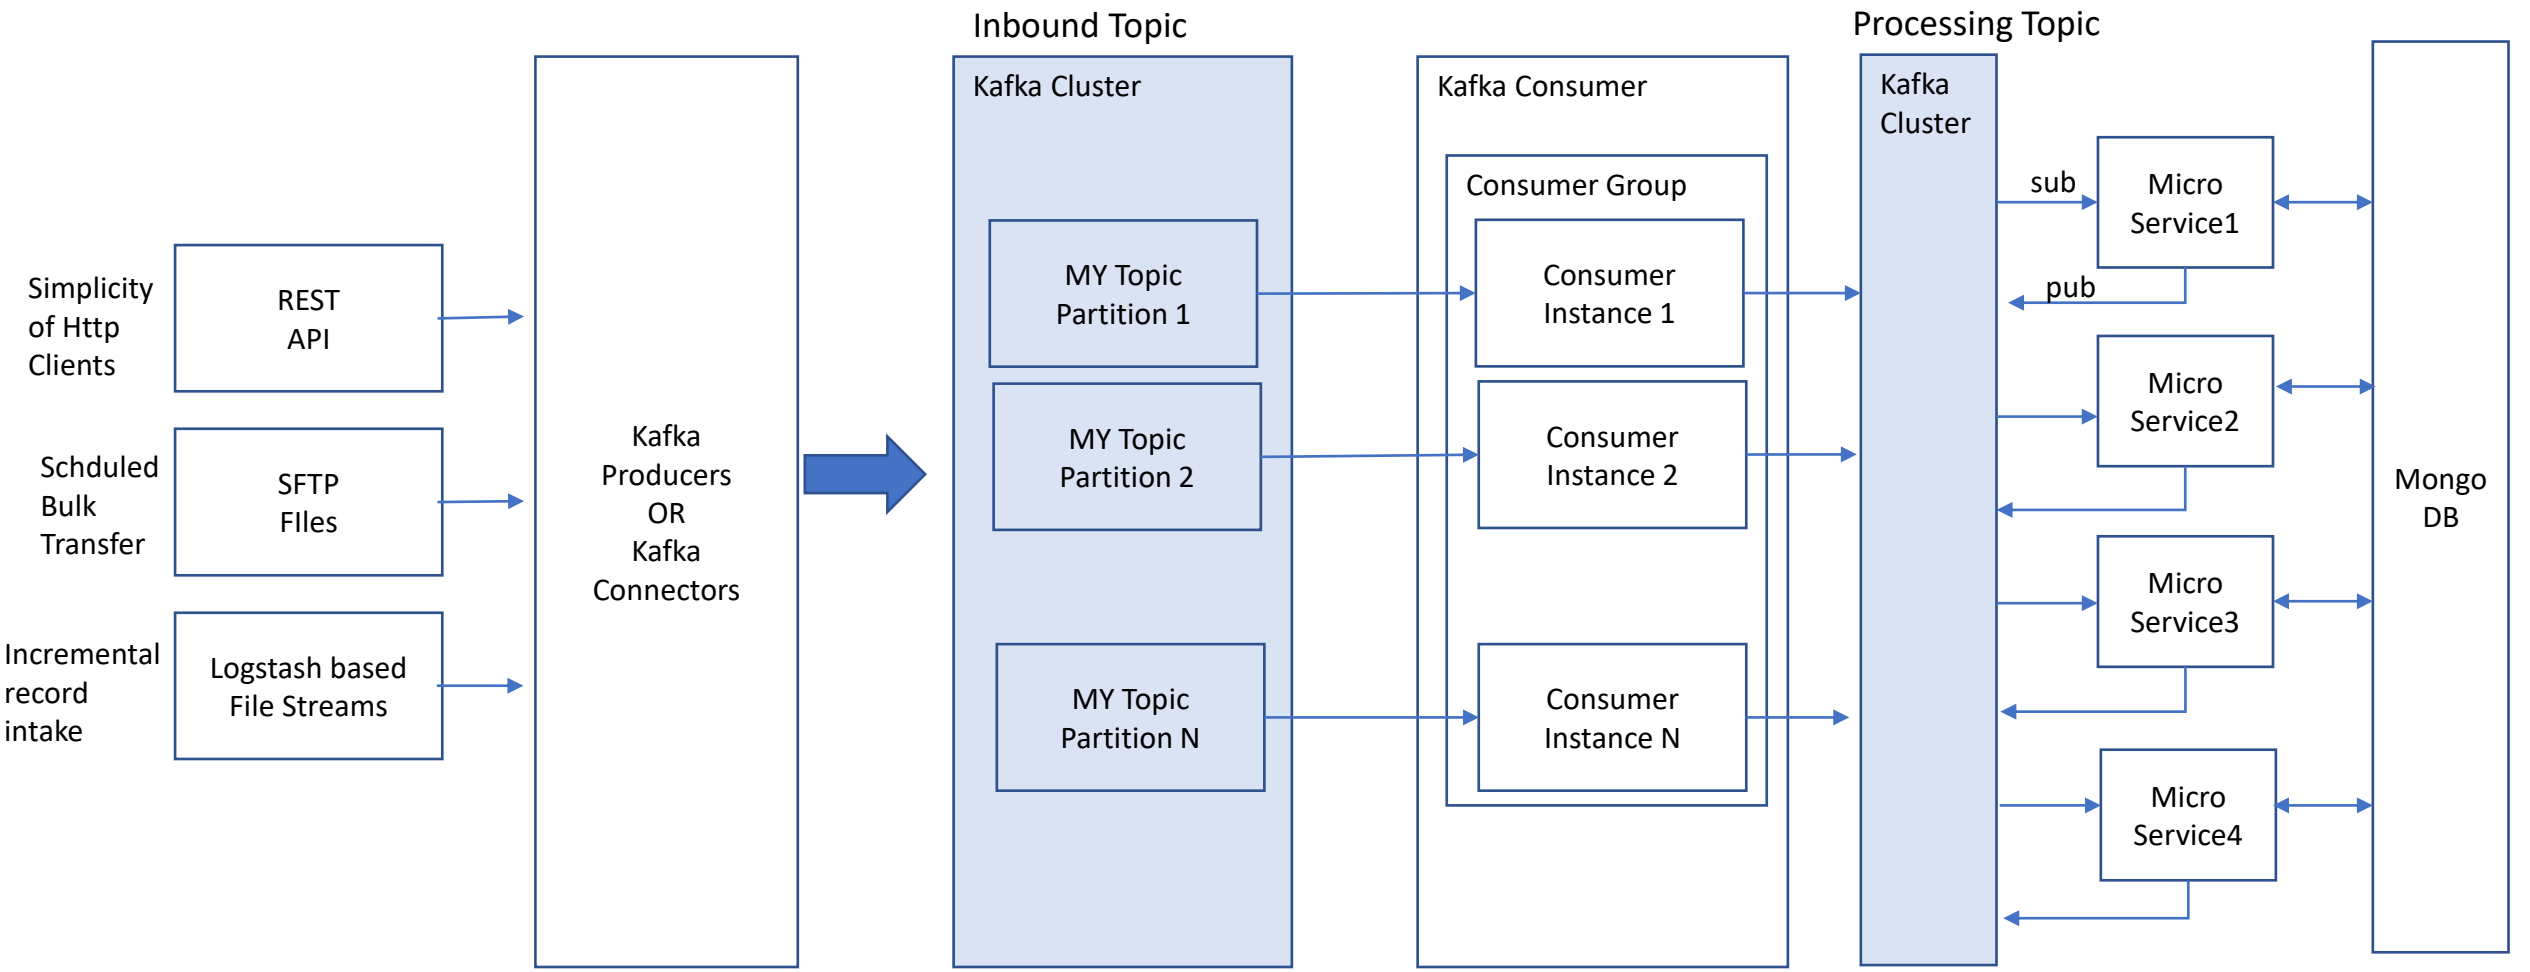

# How to decrease complexity of data model and make it more denormalized and faster to access

Org Master

| Id | Org name | Org address | Org city | Org Attr 1 | Org Attr 2 | |
|---|---|---|---|---|---|---|
| | | | | | | |

Type Master – Name Value Only

| Id | Tran Type Code | Tran Type Description |
|---|---|---|
| | | |

Transaction Table 1

| Id | Tran Type | Created on | Updated on | Tran Attr 1 | Tran Attr 2 | PartnerOrg |
|---|---|---|---|---|---|---|
| | Id OR desc | | | | | Org Id OR Org Name |

**For name value masters like Type Master:**

- Ascertain if tran type code has business meaning
- If not, code is merely a “made-up” attribute
- The code and description may never change independently

In such cases:

- Best retain Type Master Table but with single column “description” for data driven dropdowns in UI
- Store the Type description directly in the transaction table

**For classical masters like Org Master:**

- Ascertain, which Org attributes need to be shown as part of transaction, say only org name is required
- If orgname is a unique business key and cannot change over time
- Store the orgname directly in the transaction table
- Even org address can be stored in the transaction table directly
- If 2 months later, org address changes, this transaction should ideally still show the old address since that is factually correct !
- Use the Org Master only to show the Org drop downs and “current”Org Details
- No need to have foreign key contraints needlessly, between the transaction tables and masters, **if they can be avoided**